

## ▼ 1. Download the dataset

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import skew
%matplotlib inline

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.feature_selection import SelectKBest
from sklearn.metrics import r2_score, mean_squared_error
from sklearn.neighbors import KNeighborsRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.linear_model import Ridge
from sklearn.svm import SVR
import warnings
```

## ▼ 2. Load the dataset into the tool.

```
data = pd.read_csv('/content/abalone.csv')

data['age'] = data['Rings']+1.5
data.drop('Rings', axis = 1, inplace = True)

print('This dataset has {} observations with {} features.'.format(data.shape[0], data.shape[1]))

This dataset has 4177 observations with 9 features.

data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4177 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Sex             4177 non-null   object
 1   Length          4177 non-null   float64
 2   Diameter        4177 non-null   float64
 3   Height          4177 non-null   float64
 4   Whole weight    4177 non-null   float64
 5   Shucked weight  4177 non-null   float64
 6   Viscera weight  4177 non-null   float64
```

```

7   Shell weight      4177 non-null   float64
8   age               4177 non-null   float64
dtypes: float64(8), object(1)
memory usage: 293.8+ KB

```

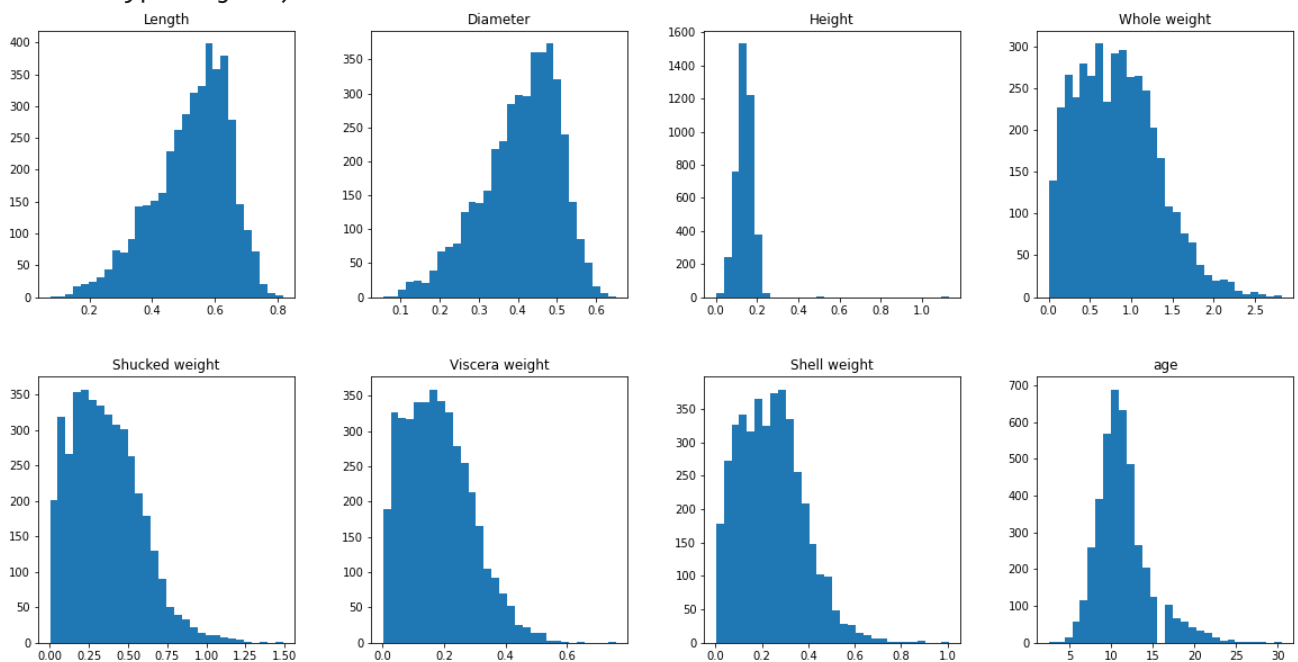
## Perform Below Visualizations: Univariate Analysis, Bi-Variate Analysis, Multi-Variate Analysis

```
data.hist(figsize=(20,10), grid=False, layout=(2, 4), bins = 30)
```

```

array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccc0d95d0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccd2a650>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccc076150>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccc02b750>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccbfe2d50>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccbf5390>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccbf5ba10>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x7f9ccbf12f90>]],
      dtype=object)

```



```
numerical_features = data.select_dtypes(include=[np.number]).columns
categorical_features = data.select_dtypes(include=[np.object]).columns
numerical_features
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2: DeprecationWarning: `
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/n
```


```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
      'Viscera weight', 'Shell weight', 'age'],
      dtype='object')
```

```
categorical_features
```

```
Index(['Sex'], dtype='object')
```

```
sns.countplot(x = 'Sex', data = data, palette="Set3")
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f9ccb5df250>
```




```
plt.figure(figsize = (20,7))
sns.swarmplot(x = 'Sex', y = 'age', data = data, hue = 'Sex')
sns.violinplot(x = 'Sex', y = 'age', data = data)
```

```

/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 56.2
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 52.2
warnings.warn(msg, UserWarning)
/usr/local/lib/python3.7/dist-packages/seaborn/categorical.py:1296: UserWarning: 58.5
warnings.warn(msg, UserWarning)
<matplotlib.axes._subplots.AxesSubplot at 0x7f9ccdf2abd0>

```



```

data.groupby('Sex')[['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
'Viscera weight', 'Shell weight', 'age']].mean().sort_values('age')

```


	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age
<b>Sex</b>								
<b>I</b>	0.427746	0.326494	0.107996	0.431363	0.191035	0.092010	0.128182	9.390462
<b>M</b>	0.561391	0.439287	0.151381	0.991459	0.432946	0.215545	0.281969	12.205497
<b>F</b>	0.579093	0.454732	0.158011	1.046532	0.446188	0.230689	0.302010	12.629304

```

sns.pairplot(data[numerical_features])


```

&lt;seaborn.axisgrid.PairGrid at 0x7f9ccbc28390&gt;

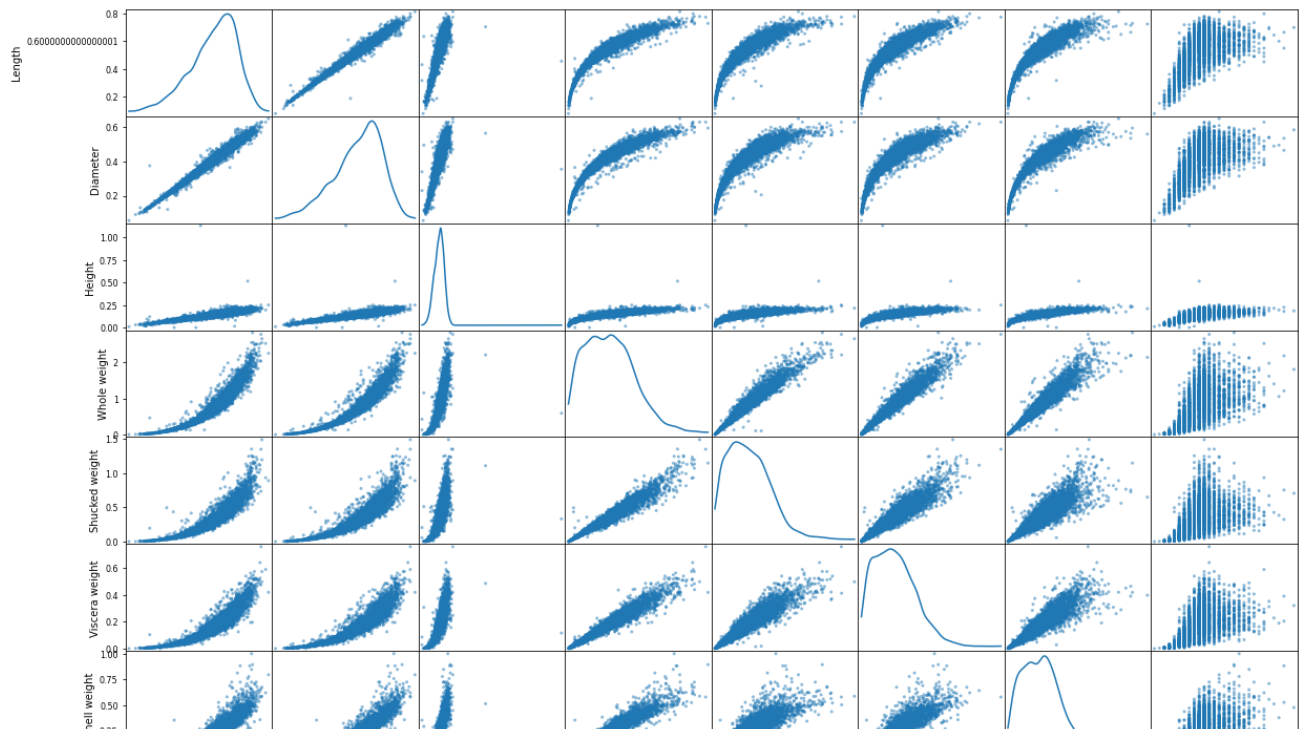


```
plt.figure(figsize=(20,7))  
sns.heatmap(data[numerical_features].corr(), annot=True)
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f9ce4c80dd0>



```
pd.plotting.scatter_matrix(data.loc[:, 'Sex':'age'], diagonal="kde",figsize=(20,15))
plt.show()
```



```
ax = data[['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',
          'Viscera weight', 'Shell weight', 'age']].plot(figsize=(20,15))
ax.legend(loc='center left', bbox_to_anchor=(1, 0.5));
```



## ▼ 4. Perform descriptive statistics on the dataset

```
data.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	
<b>count</b>	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	41
<b>mean</b>	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	
<b>std</b>	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	
<b>min</b>	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	
<b>25%</b>	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	
<b>50%</b>	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	
<b>75%</b>	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	

```
data['Sex'].describe()
```

```
count      4177
unique       3
top          M
freq       1528
Name: Sex, dtype: object
```

```
data['Sex'].value_counts()
```

```
M      1528
I      1342
F      1307
Name: Sex, dtype: int64
```

```
data.kurtosis()
```



```


/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1: FutureWarning: Droppi
    """Entry point for launching an IPython kernel.
Length          0.064621
Diameter        -0.045476
Height          76.025509
Whole weight    -0.023644
Shucked weight  0.595124
Viscera weight  0.084012
Shell weight    0.531926
age             2.330687
dtype: float64

```

```

skew_values = skew(data[numerical_features], nan_policy = 'omit')
dummy = pd.concat([pd.DataFrame(list(numerical_features), columns=['Features']),
                    pd.DataFrame(list(skew_values), columns=['Skewness degree'])], axis = 1)
dummy.sort_values(by = 'Skewness degree' , ascending = False)

```

	Features	Skewness degree	
2	Height	3.127694	
7	age	1.113702	
4	Shucked weight	0.718840	
6	Shell weight	0.620704	
5	Viscera weight	0.591640	
3	Whole weight	0.530768	
1	Diameter	-0.608979	
0	Length	-0.639643	

## ▼ 5.Check for Missing values and deal with them.

```

missing_values = data.isnull().sum().sort_values(ascending = False)
percentage_missing_values = (missing_values/len(data))*100
pd.concat([missing_values, percentage_missing_values], axis = 1, keys= ['Missing values',

```

	Missing values	% Missing
<b>Sex</b>	0	0.0
<b>Length</b>	0	0.0



data.head

```
<bound method NDFrame.head of
Shucked weight \
0      M    0.455    0.365    0.095    0.5140    0.2245
1      M    0.350    0.265    0.090    0.2255    0.0995
2      F    0.530    0.420    0.135    0.6770    0.2565
3      M    0.440    0.365    0.125    0.5160    0.2155
4      I    0.330    0.255    0.080    0.2050    0.0895
...    ..    ...    ...    ...    ...
4172   F    0.565    0.450    0.165    0.8870    0.3700
4173   M    0.590    0.440    0.135    0.9660    0.4390
4174   M    0.600    0.475    0.205    1.1760    0.5255
4175   F    0.625    0.485    0.150    1.0945    0.5310
4176   M    0.710    0.555    0.195    1.9485    0.9455


Viscera weight  Shell weight  age
0              0.1010        0.1500  16.5
1              0.0485        0.0700   8.5
2              0.1415        0.2100  10.5
3              0.1140        0.1550  11.5
4              0.0395        0.0550   8.5
...           ...          ...    ...
4172           0.2390        0.2490  12.5
4173           0.2145        0.2605  11.5
4174           0.2875        0.3080  10.5
4175           0.2610        0.2960  11.5
4176           0.3765        0.4950  13.5
```

[4177 rows x 9 columns]>


## ➤ 6. Find the outliers and replace them outliers

```
original_data = data.copy()
data = pd.get_dummies(data)
dummy_data = data.copy()
data.boxplot( rot = 90, figsize=(20,5))
```

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f9cc3f8a090>




```
var = 'Viscera weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
# outliers removal
data.drop(data[(data['Viscera weight'] > 0.5) & (data['age'] < 20)].index, inplace=True)
data.drop(data[(data['Viscera weight'] < 0.5) & (data['age'] > 25)].index, inplace=True)
```


```
var = 'Shell weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Shell weight'] > 0.6) & (data['age'] < 25)].index, inplace=True)
```


```
data.drop(data[(data['Shell weight']<0.8) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Shucked weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```




```
data.drop(data[(data['Shucked weight']>= 1) & (data['age'] < 20)].index, inplace=True)
data.drop(data[(data['Shucked weight']<1) & (data['age'] > 20)].index, inplace=True)
```

```
var = 'Whole weight'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```




```
data.drop(data[(data['Whole weight']>= 2.5) & (data['age'] < 25)].index, inplace=True)
data.drop(data[(data['Whole weight']<2.5) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Diameter'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```




```
data.drop(data[(data['Diameter']<0.1) & (data['age'] < 5)].index, inplace=True)
data.drop(data[(data['Diameter']<0.6) & (data['age'] > 25)].index, inplace=True)
data.drop(data[(data['Diameter']>=0.6) & (data['age']< 25)].index, inplace=True)
```

```
var = 'Height'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Height']>0.4) & (data['age'] < 15)].index, inplace=True)
data.drop(data[(data['Height']<0.4) & (data['age'] > 25)].index, inplace=True)
```

```
var = 'Length'
plt.scatter(x = data[var], y = data['age'],)
plt.grid(True)
```



```
data.drop(data[(data['Length']<0.1) & (data['age'] < 5)].index, inplace=True)
```

```
data.drop(data[(data['Length']<0.8) & (data['age'] > 25)].index, inplace=True)
data.drop(data[(data['Length']>=0.8) & (data['age']< 25)].index, inplace=True)
```

## ▼ 7. Check for Categorical columns and perform encoding.

```
from sklearn import preprocessing
label = preprocessing.LabelEncoder()

original_data['Sex']= label.fit_transform(original_data['Sex'])
```

```
original_data.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age
<b>0</b>	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
<b>1</b>	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
<b>2</b>	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
<b>3</b>	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
<b>4</b>	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

## ▼ 8. Split the data into dependent and independent variables.

```
X = original_data.drop('age', axis = 1)
y = original_data['age']
```

```
X
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.1500
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.0700

y

```

0      16.5
1       8.5
2      10.5
3      11.5
4       8.5
...
4172   12.5
4173   11.5
4174   10.5
4175   11.5
4176   13.5
Name: age, Length: 4177, dtype: float64

```

## 9. Scale the independent variables

```
# Normalized Y
```

```
from sklearn import preprocessing
Y=y.values.reshape(-1,1)
```

```
normalized_Y = preprocessing.normalize(Y)
```

```
print (normalized_Y)
```

```

[[1.]
 [1.]
 [1.]
 ...
 [1.]
 [1.]
 [1.]]

```

```
# Standardized Y
```

```
standard_Y = Y.copy()
```

```
from sklearn import preprocessing
```

```
ss = preprocessing.StandardScaler()
ss.fit(standard_Y)
```

```
print (standard_Y)
```

```
[[16.5]
 [ 8.5]
 [10.5]
 ...
 [10.5]
 [11.5]
 [13.5]]
```

## ▼ 10. Split the data into training and testing

```
X = data.drop('age', axis = 1)
y = data['age']
```

```
standardScale = StandardScaler()
standardScale.fit_transform(X)
```

```
selectkBest = SelectKBest()
X_new = selectkBest.fit_transform(X, y)
```

```
X_train, X_test, y_train, y_test = train_test_split(X_new, y, test_size = 0.25)
```

## ▼ 11. Build the Model

### 12. Train the Model

```
np.random.seed(10)
def rmse_cv(model, X_train, y):
    rmse = - (cross_val_score(model, X_train, y, scoring='neg_mean_squared_error', cv=5))
    return(rmse*100)
```

```
models = [LinearRegression(),
           Ridge(),
           SVR(),
           RandomForestRegressor(),
           GradientBoostingRegressor(),
           KNeighborsRegressor(n_neighbors = 4),]
```

```
names = ['LR', 'Ridge', 'svm', 'GNB', 'RF', 'GB', 'KNN']
```

```
for model,name in zip(models,names):
    score = rmse_cv(model,X_train,y_train)
    print("{} : {:.6f}, {:.4f}".format(name,score.mean(),score.std()))
```

```
LR      : 362.218315, 31.358517
Ridge   : 367.492321, 28.808135
svm     : 389.796574, 30.207738
GNB     : 363.082453, 28.130258
```



```
RF      : 351.412403, 27.731600
GB      : 403.767053, 30.014933
```

```
def modelfit(alg, dtrain, predictors, performCV=True, printFeatureImportance=True, cv_fold
#Fit the algorithm on the data
alg.fit(dtrain[predictors], dtrain['age'])

#Predict training set:
dtrain_predictions = alg.predict(dtrain[predictors])
#dtrain_predprob = alg.predict_proba(dtrain[predictors])[:,1]

#Perform cross-validation:
if performCV:
    cv_score = -cross_val_score(alg, dtrain[predictors], dtrain['age'], cv=cv_folds,
                                scoring='r2')

#Print model report:
print ("\nModel Report")
print( "RMSE : %.4g" % mean_squared_error(dtrain['age'].values, dtrain_predictions))
print( "R2 Score (Train): %f" % r2_score(dtrain['age'], dtrain_predictions))

if performCV:
    print( "CV Score : Mean - %.7g | Std - %.7g | Min - %.7g | Max - %.7g" % (np.mean(
                                                np.min(cv_


#Print Feature Importance:
if printFeatureImportance:
    feat_imp = pd.Series(alg.coef_, predictors).sort_values(ascending=False)
    plt.figure(figsize=(20,4))
    feat_imp.plot(kind='bar', title='Feature Importances')
    plt.ylabel('Feature Importance Score')
```

## ▼ 13. Test the Model

## 14. Measure the performance using Metrics

```
# Base Model
predictors = [x for x in data.columns if x not in ['age']]
lrm0 = Ridge(random_state=10)
modelfit(lrm0, data, predictors)
```

Model Report  
RMSE : 3.593  
R2 Score (Train): 0.529894  
CV Score : Mean - -0.4503433 | Std - 0.08079434 | Min - -0.514565 | Max - -0.3061263



[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 10:57 AM

